
AI:R PLATFORM



Real-time embedded AI Governance Software

AI REGULATE PTY LTD
ABN 896 898 108 48

12th September 2025

*Commercial in Confidence: This document contains
proprietary concepts.*



AI REGULATE PTY LTD



Hebba Gandolfo

*Founder &
Principal IT Consultant*

ABOUT AI REGULATE

With over two decades of hands-on delivery in enterprise governance and solution architecture, AI Regulate draws on deep experience from high-profile technology engagements with RMIT University and the Victorian State Government.


This deep technology expertise now powers a dedicated team at AI Regulate—focused on scalable solutions that enable large organisations to safely adopt AI and realise enterprise-wide benefits.

Our flagship innovation, the ai:R Platform, is a world-first, fully integrated AI Governance Engine with embedded, real-time control and reporting.

Built on novel, patent-pending systems architecture, ai:R enforces real-time, customisable governance at the point of AI use—meeting the emerging demands of responsible deployment with precision and control.

COMPLIANCE AT THE SPEED OF AI


The ai:R platform is a world-first real-time AI governance platform.



AI is moving quickly. With increased capability comes increased risk.



Complex organisations have iterated clear needs for controlled AI usage tailored to their risk appetite.



Real-time governance is essential for the effective implementation of these controls.

HOW AI:R COMPARES

Competitor	Governance Approach	Real-Time Enforcement	Deployment Model	Audit Logging	Standards Mapping	Patent Protection	Emotional Filtering	IDAM Integration
AWS (Bedrock & SageMaker Model Monitor)	Post-hoc monitoring; custom best-practice frameworks	No	AWS cloud	Yes; configurable	Partial; AWS Responsible AI Guide	No	No	AWS IAM only; no external IDAM integration
IBM Watson OpenScale	Fairness/bias/drift monitoring	No	Cloud or on-prem	Yes; monitoring logs	Supports NIST AI RMF; OECD Principles	No	No	IBM Cloud IAM; LDAP; limited SAML support
Microsoft Azure Responsible AI (ML interpretability tools)	Pipeline-level controls; explainability; fairness checks	No	Azure cloud	Yes; diagnostic logs	Supports Microsoft Responsible AI Standard	No	No	Azure Active Directory; SAML/OAuth
Google Cloud Vertex AI	Model monitoring; bias and drift alerts	No	Google cloud	Yes; model monitor logs	Supports Google AI Principles	No	No	Google Cloud IAM; limited SAML support
AI:R	Inline, token-level interception & filtering; fully customisable per agency	Yes	Hosted on client infrastructure (proxy/extension/endpoint)	Yes; immutable, exportable	Live-linked to DTA AI Technical Standard; OECD; EU AI Act; Privacy Act	Yes	Yes	Integrates with agency IDAM (SAML; OAuth; LDAP; AAD)

THE NEED FOR EMBEDDED CONTROL

- Uncontrolled AI tools risk PII leaks, bias, and policy violations
 - Post-hoc reviews and manual audits introduce latency and gap – these are not scalable for organizational-wide AI adoption
 - Research, clinical, and administrative teams need in-line enforcement to maintain ethics and compliance
 - Without user identification, malicious or inappropriate use can go undetected, leaving institutions exposed to reputational, ethical, and legal risks.
 - Traditional oversight models rely heavily on dedicated personnel, slowing deployment and inflating operational costs—especially as AI tools scale across departments.
-

FROM RISK TO RESOLUTION

ai:R is a **world-first** Patented Artificial Intelligence (AI) governance engine with embedded, in-line, real-time perimeter enforcement.

- AI:R intercepts every prompt—scanning for PII, profanity and compliance gaps—before it reaches the AI engine.
- Real-time (~Sub-5ms) enforcement ensures policies stay live, not static: rules are embedded inline, not bolted on post-hoc.
- Immutable audit trails provide essential monitoring for users
- Customisable governance levels pointed to live documents or published standards
- AI use with Identified users only, limits any malicious use and identifies instances.
- AI:R reduces cost of AI governance by automating oversight and reducing reliance on dedicated personnel.

WHERE TO NEXT?

Contact us to arrange an obligation-free consultation for your enterprise organisation.

[Ai.regulate@outlook.com](mailto:ai.regulate@outlook.com)

Enterprise scale software pilots are available for the ai:R Platform. Inquiries are welcome.

REFERENCES

Studies & Standards informing the ai:R Platform design

- CAIDE (Centre for Artificial Intelligence and Digital Ethics) <https://www.unimelb.edu.au/caide>
 - IBAC – Integrity risks associated with artificial intelligence <https://www.ibac.vic.gov.au/Integrity-risks-associated-with-artificial-intelligence>
 - Essential Eight Cyber Security Standards (ACSC) <https://www.cyber.gov.au/resources-business-and-government/essential-cybersecurity/essential-eight>
 - Profanity & PII filtering standard (Azure Content Moderator) <https://learn.microsoft.com/en-us/azure/cognitive-services/content-moderator/content-moderator-overview>.
 - TEQSA. Gen AI strategies for Australian higher education: Emerging practice. 28 November 2024. <https://www.teqsa.gov.au/guides-resources/resources/corporate-publications/gen-ai-strategies-australian-higher-education-emerging-practice>
-